

# EV Population Data Pipeline Implementation & Demo





# TOC

Overview

Pipeline Flow Architecture

Demo



# Overview

As a part of the interview process for the Snowflake SI Partner Data Cloud Architect role, the hiring team would assess and review the candidate's technical hands-on skills in Spark, data processing, and data pipeline design.

This assessment requires a demo implementation exercise, which requires the candidate to:

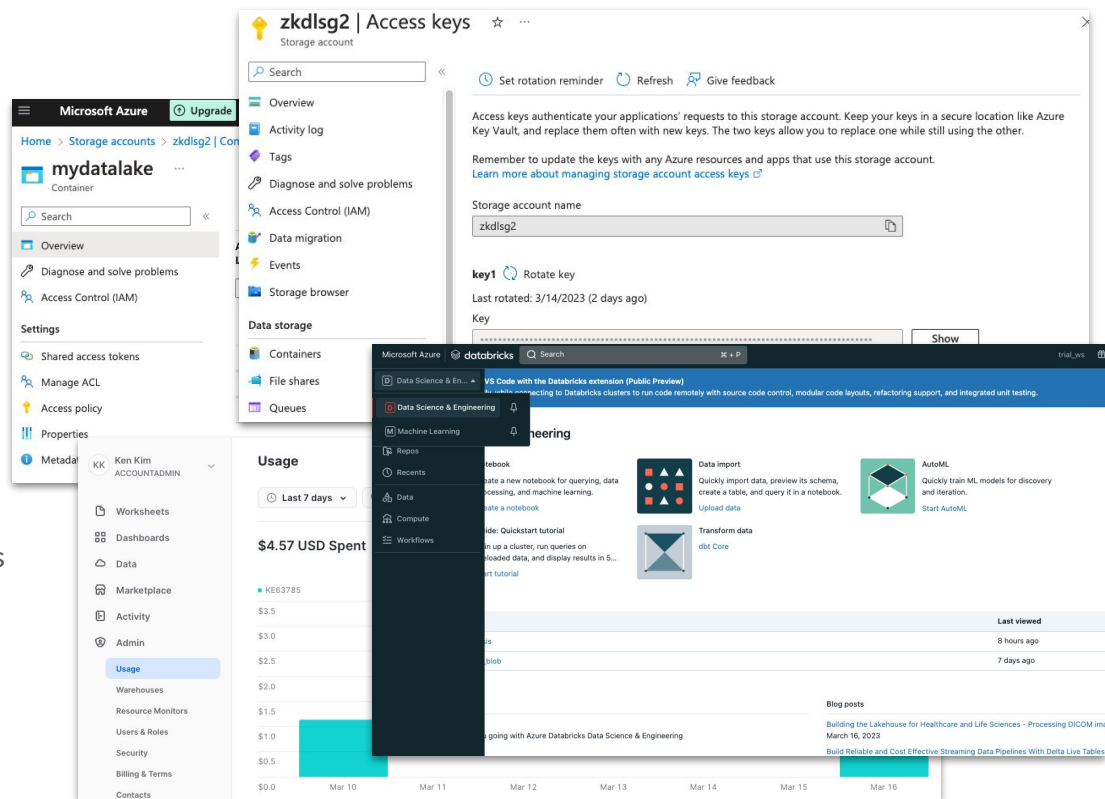
- build a data pipeline that parses the raw data from an Electric Vehicle Population Dataset (provided in a JSON file),
- extract the necessary information, and ingest it into Delta tables and then into DB tables.
- create insightful analysis based on the datasets.



# Preparation

In prior to the code implementation, below are the steps I have prepared as the prerequisite setup and due diligence checks:

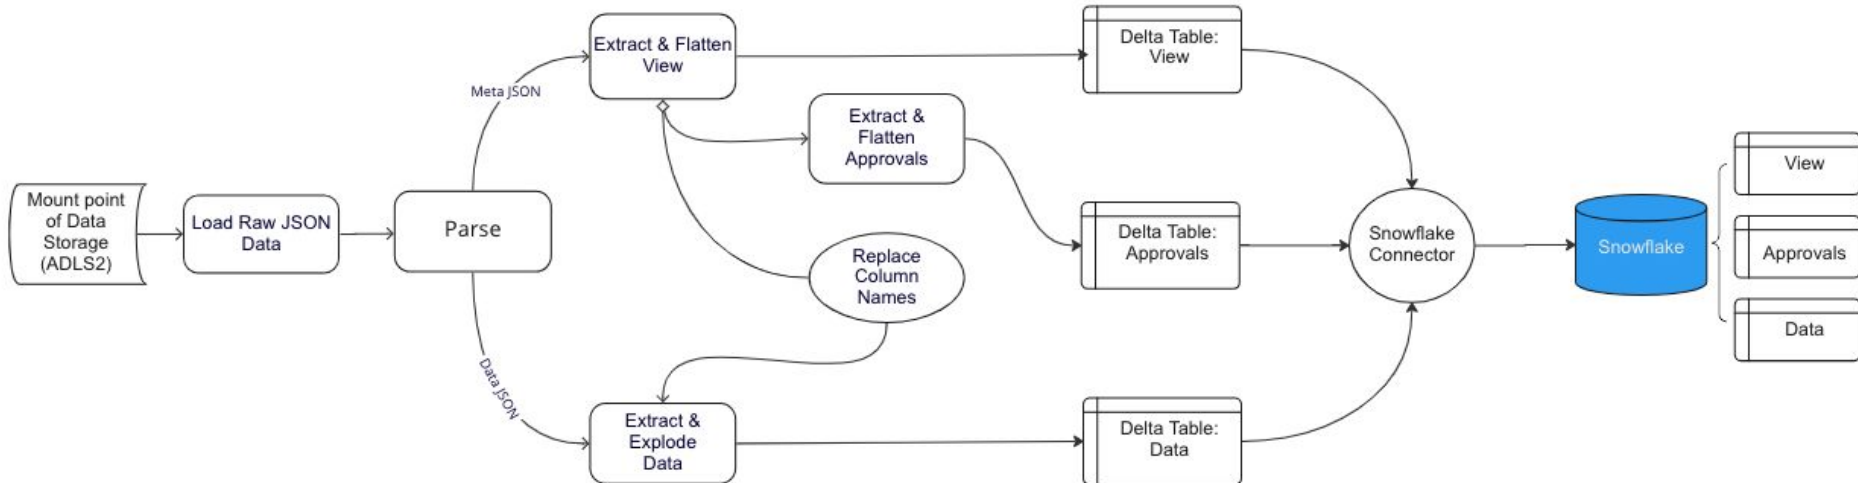
- Review of the JSON format and identify the key structures to parse and transform
- Setup Azure env (trial RG) for Storage, Access permission, etc..
- Snowflake trial sign-up
- Databricks trial sign-up
- Etc.





# Pipeline Architecture

## Data Ingestion Pipeline





# Demo

## 01: Databricks/Spark

Using Databricks Notebooks, created a Job pipeline which:

- Loads the raw data from ADLS2 storage
- Creates Spark dataframes to parse and extracts JSON objects, meta/view, meta/view/approvals, meta/view/columns and data[]
- Creates Delta Tables to flatten and store the parsed data
- Inserts the same data into Snowflake tables using Snowflake python connector

And created insights analysis example queries.

The screenshot shows a Databricks notebook interface. At the top, there's a header with 'databricks' logo, a search bar, and user information 'trial\_ws' and 'zenken@gmail.com'. Below the header, the notebook title is 'insight\_analysis' with a 'Python' language selector. The interface includes standard menu items (File, Edit, View, Run, Help) and buttons for 'Run all', 'KEN K's Cluster - sing...', 'Schedule', and 'Share'. The main area contains a Python script and its output.

**Based on the data, which car make and model would you recommend?**

```

1 #found string data type in Electric Range. fix the prob first and drop rows have any null column value and the rows that have zero-value Electric_Rnage
2 dfBase = dfData.select("Model Year", "Make", "Model", F.col("Electric Range").cast("int").alias("Electric_Range")).na.drop().filter("Electric_Range > 0")
3
4 dfficiency = dfBase.groupBy("Make", "Model")\
5   .agg(F.count("*").alias("count"), F.avg("Electric_Range").alias("Avg_Electric_Range"))\
6   .orderBy("Avg_Electric_Range", ascending=False).limit(20)
7 dfficiency.show()
8
9 dfficiency2 = dfBase.groupBy("Model Year", "Make", "Model")\
10  .agg(F.count("*").alias("count"), F.avg("Electric_Range").alias("Avg_Electric_Range"))\
11  .orderBy("Avg_Electric_Range", ascending=False).limit(20)
12 dfficiency2.show()

```

**Output:**

(4) Spark Jobs

- dfBase: pyspark.sql.dataframe.DataFrame = [Model Year: string, Make: string ... 2 more fields]
- dfficiency: pyspark.sql.dataframe.DataFrame = [Make: string, Model: string ... 2 more fields]
- dfficiency2: pyspark.sql.dataframe.DataFrame = [Model Year: string, Make: string ... 3 more fields]

Make	Model	count	Avg_Electric_Range
TESLA	MODEL Y	433	291.0
HYUNDAI	KONA	43	258.0
CHEVROLET	BOLT EV	773	243.27037516170762
TESLA	MODEL X	575	241.47130434782608
TESLA	ROADSTER	7	237.85714285714286
TESLA	MODEL 3	2541	237.61668634395906
JAGUAR	I-PACE	35	234.0
POLESTAR	PS2	9	233.0
TESLA	MODEL S	1147	226.6713164777681
AUDI	E-TRON SPORTBACK	13	218.0
AUDI	E-TRON	101	208.2722277227227
PORSCHE	TAYCAN	31	198.83225806451613
KIA	NIRO	316	143.28481012658227
MINI	HARDTOP	17	110.0
VOLKSWAGEN	E-GOLF	194	108.97938144329896
TOYOTA	RAV4	8	103.0
NISSAN	LEAF	2199	102.13278763074125
THINKI	CTTY	1	100.0

Command took 1.78 seconds -- by zenken@gmail.com at 3/16/2023, 12:09:22 AM on KEN K's Cluster - single node



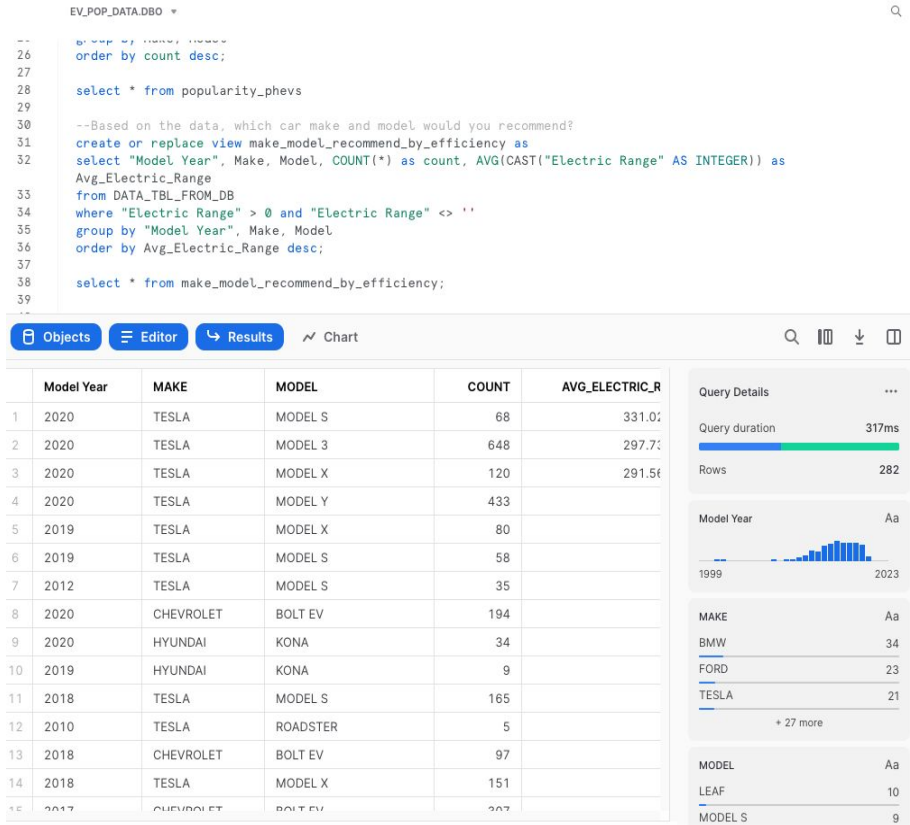
# Demo

## 02: Snowflake/Snowpark

Having the same extracted data loaded into Snowflake tables, generated view for the same insight analysis queries.

Created a Snowpipe to process the same JSON data ingestion to compare to Databricks Job Workflow.

Used Snowpark API to perform the same Dataframe





# Snowpark vs Databricks Spark

Snowpark and Databricks Spark are both powerful tools for data processing and analysis. However, there are several benefits to using Snowpark over Databricks Spark:

**Performance:** Snowpark is designed to optimize performance for Snowflake's data warehouse environment. This means that Snowpark jobs can run faster and more efficiently than Databricks Spark jobs when working with Snowflake.

**Integration with Snowflake:** Snowpark is tightly integrated with Snowflake, which makes it easy to write Spark code that can interact with Snowflake data warehouses. This means that you can use Snowpark to load data into Snowflake, extract data from Snowflake, and run complex analytical queries on Snowflake data.

**Flexibility:** Snowpark is a more flexible tool than Databricks Spark, as it can be used in a variety of environments, including on-premises data centers and public clouds. This means that you can use Snowpark to build data processing pipelines that meet your specific business needs, regardless of your organization's IT infrastructure.

Overall, Snowpark offers several benefits over Databricks Spark when working with Snowflake data warehouses. If your organization is already using Snowflake and needs a high-performance, flexible tool for data processing and analysis, Snowpark is definitely worth considering.





# Thank you.

